

Voxel-Based Screen Space Ambient Occlusion

1. INTRODUCTION

One of the biggest challenges in photo-realistic rendering is calculating realistic lighting and shadowing effects. Real world lighting is created by the interaction of photons with objects. We perceive an object when photons emitted from a light source hit an object, bounce off, then reach our eyes. Ray-tracing methods most closely approximate this by tracing the path from each point in the scene back to the camera and integrating all contributions over the volume of the scene. This results in very realistic-looking scenes, but it is very time consuming. In real-time computer graphics, lights are generally modeled either as directional or ambient. Directional lighting travels in one direction and originates at some source that may be within the field of view or a distant point outside of the field of view. Ambient light accounts for the many times that light bounces off different objects so that objects that do not directly face the light source are still illuminated. Ambient lighting is generally modeled as isotropic lighting that originates at every point in space. For directional lighting, shadows are used to darken the areas that are blocked from interaction with the light. When objects are close to each other, they can also block some of the ambient light. Ambient occlusion accounts for this effect, resulting in a much more realistic scene.

The basic principle behind ambient occlusion methods is to look at each point and determine what fraction of the surrounding volume is occupied. The point is then darkened by an amount proportional to this value. Computing this volumetric integral can be rather complicated; however, we believe voxelization can simplify this calculation. Voxelization involves rendering a scene by breaking all of the objects up into volumetric elements, called voxels, rather than using triangle meshes to approximate the shape of an object. Approximating the integral as a sum of these elements should simplify the volume calculations, resulting in a faster overall rendering of the scene.

2. RELATED WORK

Ambient occlusion techniques are grouped into two types: object-space methods, which operate over all points, and screen-space methods, which operate only over points that are visible in the final displayed image. For complex scenes, object-space methods are very expensive and are not practical for real-time applications. The first implementation of screen-space ambient occlusion was in *Crysis* [1]. In this implementation, the Z-buffer was the only input needed to calculate the ambient occlusion. For each pixel in the scene, surrounding pixels are tested to determine whether they are in front of the pixel in question. The occlusion values are then weighted by their distance from the pixel in question, then added to produce the final ambient occlusion value.

McGuire [2] computes what he calls "ambient occlusion volumes" by forming a bounding volume of polygons around the pixel and combining the contributions of all fragments within this volume using an integral approximation developed by Baum, et al. [3]. McGuire achieves visually appealing results with his algorithm, but the computation cannot be accomplished in real-time. The work most relevant to our project is the work by Penmatsa and Wyman [4]. They compute ambient occlusion using voxels and sum up the contribution of each voxel face. They also measure the curvature of objects and use this for multi-resolution effects, which they implement using mipmaps. Their solution runs at 40fps, but the quality of the images diverges considerably from ground truth.

3. ANALYTIC VOLUMETRIC OCCLUSION

Let X be an infinitesimal smooth patch, and without loss of generality let the center be at the origin with normal \hat{n} , as shown in Figure 1. Ambient illumination is the incident light on X , due to objects outside of the patch. Ambient occlusion over a hemisphere due to a voxel is computationally expensive since individual sides of voxels need to be projected onto the hemisphere to calculate total contribution per patch X [4]. Instead, we derive an ambient occlusion equation from the area representation of the rendering equation. We feel this is a reasonable assumption as the voxel volume is relatively small. Furthermore, ambient occlusion is a low frequency effect so a coarse geometry representation suffices for this computation. Let P be a voxel with vertices $\{\hat{p}_0, \hat{p}_1, \dots, \hat{p}_7\}$, and let \hat{c}_0 be the centroid of the voxel with normal \hat{N} . The occlusion by P of ambient light directed to X is equal to the form factor that describes the

diffuse radiative transfer between P and X as shown below:

$$AO_P(\hat{n}) = \frac{\rho}{2\pi} \frac{\cos(\theta_P) \cdot \cos(\theta_X)}{r^2} A_P$$

where $\cos(\theta_P) = \hat{n} \cdot \psi$, $\cos(\theta_X) = \hat{N} \cdot \psi$, ψ is the direction of the incoming radiance, A_P denotes the screen projected volume of the voxel, and r^2 represents the squared distance between c_0 and X. Lastly, ρ is a normalizing constant for ambient occlusion values. We may interpret $\cos(\theta_P) \cdot \cos(\theta_X)$ as the amount of radiance reaching point X depending on the angle made by ψ with the normals of c_0 and X. The corrected integral approximation was first introduced by Baum et al. in the context of the radiosity algorithm [3].




FIGURE 1. Diagram of sample patch X being occluded by voxel P

4. AMBIENT OCCLUSION VOXEL ALGORITHM

This project aims to introduce a new approximation algorithm for the voxel-based screen space ambient occlusion problem. It combines previous work to achieve substantially improved quality over fast methods and substantially improved performance compared to more accurate methods. The implementation of computing the analytic solution for ambient occlusion was written as a CUDA program while voxelization was implemented using an existing binary voxelization algorithm, bin-vox [5]. The main contribution of this project is an efficient algorithm for estimating ambient occlusion using an analytic solution to the integral form of the full matrix radiosity method as described in the previous section.




FIGURE 2. Binary voxelized bunny.obj scene using $256 \times 256 \times 256$ dicing planes in each direction. [5]

Preliminary ambient occlusion algorithm.

- (1) Render voxelized scene
- (2) For each screen pixel:
 - Unproject to world space coordinates and then project to voxel space coordinates
 - Determine the position and normal of the centroid of voxels
 - Determine the neighborhood of occupied voxels and their relative distance
 - Use analytic solution to compute occlusion
 - Update occlusion buffer accordingly
- (3) Render the normalized occlusion buffer

On a first pass, we render the scene using OpenGL and output the list of screen pixels from the current camera view. Voxels are generated on a second using binvox [5], which reads in a 3D model file, rasterizes it into a binary 3D uniform voxel grid, and writes the resulting voxel file. We would like to extend this work and render the voxelized scene only from the viewing frustum as this would be more efficient than the current method. We found that the CUDA 4.0 layered texture memory access was a natural fit to our description of the binary voxelized scene. This decision allowed us to take advantage of the 2D spatially arranged texture cache which guarantees scalability of our algorithm.

A list of camera-accessed screen pixels is a good representation for X ; however, for our computation we need to unproject a screen pixel to object space and then obtain its corresponding occupied voxel coordinates. Once we obtain the occupied voxel coordinates for the corresponding screen pixels, we compute ambient occlusion only for that list of points. Because we unproject into voxel space for the ambient occlusion calculations, sections of the object immediately outside the field of view will also be used in the occlusion calculations, which prevents the points at the window edges from appearing lighter than the rest of the scene.

We implemented both a CPU and GPU implementation of the ambient occlusion algorithm to verify our results. Given an input point from the list of voxelized screen pixels, we compute its neighbors within a 4 voxel-wide cube and compute their distance weighted by the presence value (0 indicates an empty voxel and 1 is a filled voxel). We also can compute voxel normals using the gradient volume and the direction of incoming radiance between the neighboring voxel and the occluded point. Lastly, on the GPU implementation we use a reduction algorithm to sum up the preliminary ambient occlusion values per occluded point.

Currently, our GPU implementation permits two options for work allocation and scheduling. For larger data sets, we scale down the width of the neighboring cube and map each projected screen pixel to a thread on the GPU. For smaller data sets, we can expand our neighboring cube and map each projected screen pixel to a thread block, where each thread is responsible for computing ambient occlusion for a neighboring voxel.

5. RESULTS AND CONCLUSIONS

We are still in the process of completing a full implementation of the ambient occlusion algorithm. The parts left to implement are: calculating normals of the points chosen from screen space and obtaining an accurate projection of occlusion data in the scene. Because we do not have the object normals, we are approximating the ambient occlusion value using the sum of the square distances from X to its occluding voxels. This gives us a value that is inaccurate, but should nevertheless result in darkening of a point proportional to the number and distance of occluding voxels. For our tests we used two different data sets: one was the list we generated containing all filled voxels visible in screen space; the other was the complete list of filled voxels in our three dimensional test model. We timed the GPU processing time (without accounting for

setup in texture memory and voxelization) for both data sets. The run-times are 0.303 ms for the screen space projected data set, and 7.016 ms for the full voxel space data set. The calculation for the reduced data set is over 23 \times faster because most voxels in the model are not visible in a single frame. This demonstrates the benefit of screen space ambient occlusion methods over object space methods. One interesting observation was that the calculation for the complete data set used only 2391 million texture lookups per second (Mtex/s), while the reduced screen space data set required 55370 Mtex/s. This is because the full data set benefits more from locality than the sparse screen space data set.



FIGURE 3. Ambient occlusion calculated using NVIDIA's OptiX ray-tracer.

Converting the voxel coordinates back into object space, then screen space yielded Figure 4 and Figure 5, which show the darkening calculated by our implementation. These images are obviously not correct, but they do reveal a few interesting points. As expected, the scene computed with the full voxel space description is much darker than the scene generated using only screen space values. This is because the shadowing that should be blocked by other parts of the object is not eliminated.

One aspect of these results that seems hopeful is that the darkness does seem to have the general shape of a bunny, as shown in the up close view in Figure 6. This indicates that there is probably an error in the scaling calculations that needs to be identified. Also, the occlusion tends to be grouped into smaller blocks on the screen. This probably

comes from the fact that, using our current methods, we are projecting all points within one voxel back to the same window coordinate. This problem could be solved by keeping a list of all pixel-voxel pairs in the scene, rather than projecting occlusion straight from voxel coordinates. Lastly, our code is available on Github at [git://github.com/beezees/ambientocclusion.git](https://github.com/beezees/ambientocclusion.git).



FIGURE 4. Closer view of full voxel space ambient occlusion from Figure 6.




FIGURE 5. Ambient occlusion values calculated using screen space values only.

REFERENCES

- [1] Mittring, M.: "Finding next gen: Cryengine 2". In: ACM SIGGRAPH 2007 Courses, pp. 97–121 (2007).
- [2] McGuire, M.: "Ambient Occlusion Volumes". In: Proceedings of High Performance Graphics, pp. 47-56 (2010).
- [3] Baum, D., Rushmeier, H., Winget, J.: "Improving Radiosity Solutions Through the Use of Analytically Determined Form-Factors." Computer Graphics 23(3), pp. 325–334 (1989).
- [4] Penmatsa, R., Wyman, C.: "Voxel-Space Ambient Occlusion". (UICS-12-01) Iowa City: The University of Iowa (2012).
- [5] Turk G., Nooruddin F.: "Simplification and Repair of Polygonal Models Using Volumetric Techniques". Gvu technical report 99-37 (later published in IEEE Trans. on Visualization and Computer Graphics, vol. 9, nr. 2, April 2003, pages 191-205).




FIGURE 6. Ambient occlusion values calculated using voxel coordinates of the full model.